

Elements Of Data Science - F2025

Week 8: Data Cleaning and Feature Engineering

11/11/2025

Homework 3 due on 11/25/2025 @ 11:59 pm


Weekly Quiz due on Nov 18 @ 11:59 pm

TODOs

- Readings:
 - PML Ch4.5 : Selecting Meaningful Features
 - PML Ch5.1 : Unsupervised dimensionality reduction via principal component analysis
 - [Recommended] Pandas: Merge, join, concatenate and compare
 - [Additional] PDSH 5.9 : PCA
 - [Optional] : Nice ROC visualization (<http://www.navan.name/roc/>)

Precision & Recall and ROC visualizations

Precision & Recall




How many selected items are relevant?

$$\text{Precision} = \frac{\text{true positives}}{\text{true positives} + \text{false positives}}$$

How many relevant items are selected?

$$\text{Recall} = \frac{\text{true positives}}{\text{true positives} + \text{false negatives}}$$

ROC




This and more at https://github.com/dariyasydykova/open_projects/tree/master/ROC_animation

Also see the interactive viz at <http://www.navan.name/roc/>

Precision & Recall and ROC visualizations

Precision & Recall




How many selected items are relevant?

$$\text{Precision} = \frac{\text{true positives}}{\text{true positives} + \text{false positives}}$$

How many relevant items are selected?

$$\text{Recall} = \frac{\text{true positives}}{\text{true positives} + \text{false negatives}}$$

ROC



This and more at https://github.com/dariyasydykova/open_projects/tree/master/ROC_animation

Also see the interactive viz at <http://www.navan.name/roc/>

Today

- Data Cleaning
 - Duplicates
 - Missing Data
 - Dummy Variables
 - Rescaling
 - Dealing With Skew
 - Removing Outliers
- Feature Engineering
 - Binning
 - One-Hot encoding
 - Derived
 - string functions
 - datetime functions

Questions?

Environment Setup

Environment Setup

In [99]:

```
1 import numpy
2 import numpy as np
3 import pandas as pd
4 import matplotlib.pyplot as plt
5 import seaborn as sns
6
7 from mlxtend.plotting import plot_decision_regions
8
9 sns.set_style('darkgrid')
10
11 %matplotlib inline
```

Data Cleaning

Why do we need clean data?

- **To remove duplicates** Want one row per observation
- **To remove/fill missing** Most models cannot handle missing data
- **To engineer features** Most models require fixed length feature vectors
- Different models require different types of data (transformation)
 - **Linear models:** real valued features with similar scale
 - **Distance based:** real valued features with similar scale
 - **Tree based:** can handle unscaled real and categorical

Example Data

Example Data

In [100]:

```
1 # read in example data
2 df_shop_raw = pd.read_csv('../data/flowershop_data_with_dups_week8.csv',
3                           header=0,
4                           delimiter=',')
5 df_shop_raw['purchase_date'] = pd.to_datetime(df_shop_raw.purchase_date)
6
7 # make a copy for editing
8 df_shop = df_shop_raw.copy()
9
10 df_shop.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1001 entries, 0 to 1000
Data columns (total 6 columns):
 #   Column           Non-Null Count  Dtype  
---  --  
 0   purchase_id     1001 non-null    int64  
 1   lastname        1001 non-null    object  
 2   purchase_date   1001 non-null    datetime64[ns]
 3   stars           1001 non-null    int64  
 4   price           979 non-null    float64 
 5   favorite_flower 823 non-null    object  
dtypes: datetime64[ns](1), float64(1), int64(2), object(2)
memory usage: 47.0+ KB
```

Duplicated Data

- Only drop duplicates if you know data should be unique
 - Example: if there is a unique id per row

Duplicated Data

- Only drop duplicates if you know data should be unique
 - Example: if there is a unique id per row

```
In [101]: 1 df_shop.duplicated().iloc[-3:] # are any of the last 3 rows duplicates?  
Out[101]: 998    False  
         999    False  
        1000   True  
       dtype: bool
```

Duplicated Data

- Only drop duplicates if you know data should be unique
 - Example: if there is a unique id per row

```
In [101]: 1 df_shop.duplicated().iloc[-3:] # are any of the last 3 rows duplicates?
```

```
Out[101]: 998    False
          999    False
         1000   True
dtype: bool
```

```
In [102]: 1 df_shop[df_shop.duplicated(keep='first')] # default: keep 'first' duplicated row
```

```
Out[102]:
purchase_id  lastname  purchase_date  stars  price  favorite_flower
1000      1010     FERGUSON  2017-05-04     2    21.02    daffodil
```

Duplicated Data

- Only drop duplicates if you know data should be unique
 - Example: if there is a unique id per row

```
In [101]: 1 df_shop.duplicated().iloc[-3:] # are any of the last 3 rows duplicates?
```

```
Out[101]: 998    False
999    False
1000   True
dtype: bool
```

```
In [102]: 1 df_shop[df_shop.duplicated(keep='first')] # default: keep 'first' duplicated row
```

```
Out[102]:
purchase_id  lastname  purchase_date  stars  price  favorite_flower
1000  1010      FERGUSON  2017-05-04    2    21.02  daffodil
```

```
In [103]: 1 df_shop[df_shop.duplicated(keep=False)] # keep=False to show all duplicated rows
```

```
Out[103]:
purchase_id  lastname  purchase_date  stars  price  favorite_flower
10    1010      FERGUSON  2017-05-04    2    21.02  daffodil
1000  1010      FERGUSON  2017-05-04    2    21.02  daffodil
```

Duplicated Data for Subset of Columns

Duplicated Data for Subset of Columns

```
In [110]: 1 # if it's important that a subset of columns is not duplicated
2 (
3     df_shop
4     .sort_values(by='purchase_id')
5     .loc[df_shop.duplicated(subset=['purchase_id']), keep =False])
6
7 )
```

Out[110]:

		purchase_id	lastname	purchase_date	stars	price	favorite_flower
	1000	1010	FERGUSON	2017-05-04	2	21.02	daffodil
	10	1010	FERGUSON	2017-05-04	2	21.02	daffodil
	100	1101	WEBB	2017-07-13	2	8.00	iris
	101	1101	BURKE	2017-08-16	4	18.56	daffodil

```
In [107]: 1 df_shop.loc[df_shop.duplicated(subset=['purchase_id']),:]
```

Out[107]:

		purchase_id	lastname	purchase_date	stars	price	favorite_flower
	101	1101	BURKE	2017-08-16	4	18.56	daffodil
	1000	1010	FERGUSON	2017-05-04	2	21.02	daffodil

Duplicated Data for Subset of Columns

```
In [110]: 1 # if it's important that a subset of columns is not duplicated
2 (
3     df_shop
4     .sort_values(by='purchase_id')
5     .loc[df_shop.duplicated(subset=['purchase_id']), keep =False])
6
7 )
```

Out[110]:

		purchase_id	lastname	purchase_date	stars	price	favorite_flower
	1000	1010	FERGUSON	2017-05-04	2	21.02	daffodil
	10	1010	FERGUSON	2017-05-04	2	21.02	daffodil
	100	1101	WEBB	2017-07-13	2	8.00	iris
	101	1101	BURKE	2017-08-16	4	18.56	daffodil

```
In [107]: 1 df_shop.loc[df_shop.duplicated(subset=['purchase_id']),:]
```

Out[107]:

		purchase_id	lastname	purchase_date	stars	price	favorite_flower
	101	1101	BURKE	2017-08-16	4	18.56	daffodil
	1000	1010	FERGUSON	2017-05-04	2	21.02	daffodil

```
In [105]: 1 # could also use the indexing shortcut
2 df_shop[df_shop.duplicated(subset=['purchase_id'],keep='first')].sort_values(by='purchase_id')
```

Out[105]:

		purchase_id	lastname	purchase_date	stars	price	favorite_flower
	1000	1010	FERGUSON	2017-05-04	2	21.02	daffodil
	101	1101	BURKE	2017-08-16	4	18.56	daffodil

Dropping Duplicates

Dropping Duplicates

```
In [9]: 1 df_shop = df_shop.drop_duplicates(subset=None,      # consider all columns  
2                               keep='first',      # or 'last' or False)  
3                               inplace=False     # default, return a modified dataframe  
4 )
```

Dropping Duplicates

In [9]:

```
1 df_shop = df_shop.drop_duplicates(subset=None,      # consider all columns
2                                   keep='first',    # or 'last' or False)
3                                   inplace=False)   # default, return a modified dataframe
4
```

In [10]:

```
1 # or can use inplace to change the original dataframe
2 df_shop.drop_duplicates(subset=None,
3                        keep='first',
4                        inplace=True) # change the dataframe in place
5
```

Dropping Duplicates

```
In [9]: 1 df_shop = df_shop.drop_duplicates(subset=None,      # consider all columns  
2                                         keep='first',    # or 'last' or False)  
3                                         inplace=False     # default, return a modified dataframe  
4 )
```

```
In [10]: 1 # or can use inplace to change the original dataframe  
2 df_shop.drop_duplicates(subset=None,  
3                         keep='first',  
4                         inplace=True # change the dataframe in place  
5 )
```

```
In [11]: 1 # can also drop rows with duplicates considering only a subset of columns  
2 df_shop = df_shop.drop_duplicates(subset=['purchase_id'])
```

Missing Data

- Reasons for missing data
 - Sensor error (random?)
 - Data entry error (random?)
 - Survey-subject decisions (non-random?)
 - etc.

Missing Data

- Reasons for missing data
 - Sensor error (random?)
 - Data entry error (random?)
 - Survey-subject decisions (non-random?)
 - etc.
- Dealing with missing data
 - Drop rows
 - Impute from data in the same column
 - Infer from other features
 - Fill with adjacent data

Missing Data in Pandas: `np.nan`

- Missing values represented by `np.nan` : Not A Number

Missing Data in Pandas: np.nan

- Missing values represented by np.nan : Not A Number

```
In [12]: 1 # Earlier, we saw missing values in the dataframe summary  
2 # df_shop.info()
```

Missing Data in Pandas: np.nan

- Missing values represented by np.nan : Not A Number

```
In [12]: 1 # Earlier, we saw missing values in the dataframe summary  
2 # df_shop.info()
```

```
In [13]: 1 # can we check for NaN using "x == np.nan"? No!  
2 np.nan == np.nan  
3
```

```
Out[13]: False
```

Missing Data in Pandas: np.nan

- Missing values represented by np.nan : Not A Number

```
In [12]: 1 # Earlier, we saw missing values in the dataframe summary  
2 # df_shop.info()
```

```
In [13]: 1 # can we check for NaN using "x == np.nan"? No!  
2 np.nan == np.nan  
3
```

```
Out[13]: False
```

`==` operator compares the values of both the operands and checks for value equality. `is` operator, on the other hand, checks whether both the operands refer to the same object or not.

IEEE 754 Standard: The standard specifically dictates that NaN should not be equal to anything, including itself. This design choice is based on the idea that if you have two unknown or invalid values, you cannot definitively say they are the same.

Missing Data in Pandas: np.nan

- Missing values represented by np.nan : Not A Number

```
In [12]: 1 # Earlier, we saw missing values in the dataframe summary  
2 # df_shop.info()
```

```
In [13]: 1 # can we check for NaN using "x == np.nan"? No!  
2 np.nan == np.nan  
3
```

```
Out[13]: False
```

`==` operator compares the values of both the operands and checks for value equality. `is` operator, on the other hand, checks whether both the operands refer to the same object or not.

IEEE 754 Standard: The standard specifically dictates that NaN should not be equal to anything, including itself. This design choice is based on the idea that if you have two unknown or invalid values, you cannot definitively say they are the same.

```
In [14]: 1 # however  
2 np.nan is np.nan
```

How to check for NaN: `.isna()` and `.notna()`

How to check for NaN: `.isna()` and `.notna()`

```
In [18]: 1 # some missing data
          2 df_shop.loc[20:21, 'price']
```

```
Out[18]: 20      NaN
          21    10.53
Name: price, dtype: float64
```

How to check for NaN: `.isna()` and `.notna()`

```
In [18]: 1 # some missing data  
2 df_shop.loc[20:21, 'price']
```

```
Out[18]: 20      NaN  
21    10.53  
Name: price, dtype: float64
```

```
In [19]: 1 # .isna() returns True where data is missing, False otherwise  
2 df_shop.loc[20:21, 'price'].isna()
```

```
Out[19]: 20    True  
21   False  
Name: price, dtype: bool
```

How to check for NaN: `.isna()` and `.notna()`

```
In [18]: 1 # some missing data  
2 df_shop.loc[20:21, 'price']
```

```
Out[18]: 20      NaN  
21    10.53  
Name: price, dtype: float64
```

```
In [19]: 1 # .isna() returns True where data is missing, False otherwise  
2 df_shop.loc[20:21, 'price'].isna()
```

```
Out[19]: 20    True  
21   False  
Name: price, dtype: bool
```

```
In [20]: 1 # .notna() returns True where data is NOT missing, False otherwise  
2 df_shop.loc[20:21, 'price'].notna()
```

```
Out[20]: 20   False  
21    True  
Name: price, dtype: bool
```

How to check for NaN: `.isna()` and `.notna()`

In [18]:

```
1 # some missing data
2 df_shop.loc[20:21, 'price']
```

Out[18]:

```
20      NaN
21    10.53
Name: price, dtype: float64
```

In [19]:

```
1 # .isna() returns True where data is missing, False otherwise
2 df_shop.loc[20:21, 'price'].isna()
```

Out[19]:

```
20    True
21   False
Name: price, dtype: bool
```

In [20]:

```
1 # .notna() returns True where data is NOT missing, False otherwise
2 df_shop.loc[20:21, 'price'].notna()
```

Out[20]:

```
20   False
21   True
Name: price, dtype: bool
```

In [21]:

```
1 # find rows where price is missing
2 df_shop[df_shop.price.isna()].head()
```

Out[21]:

	purchase_id	lastname	purchase_date	stars	price	favorite_flower
20	1020	CLARK	2017-01-05	3	NaN	NaN
41	1041	PETERS	2017-02-01	4	NaN	orchid
54	1054	GREEN	2017-02-13	5	NaN	daffodil
63	1063	BARNETT	2017-08-27	4	NaN	gardenia
145	1145	CARROLL	2017-07-29	3	NaN	tulip

Counting NaNs

Counting NaNs

```
In [22]: 1 # How many nan's in a single column?  
2 df_shop.price.isna().sum()
```

```
Out[22]: 22
```

Counting NaNs

```
In [22]: 1 # How many nan's in a single column?  
2 df_shop.price.isna().sum()
```

```
Out[22]: 22
```

```
In [23]: 1 # How many nan's per column?  
2 df_shop.isna().sum()
```

```
Out[23]: purchase_id      0  
lastname            0  
purchase_date       0  
stars              0  
price             22  
favorite_flower    178  
dtype: int64
```

Counting NaNs

```
In [22]: 1 # How many nan's in a single column?  
2 df_shop.price.isna().sum()
```

```
Out[22]: 22
```

```
In [23]: 1 # How many nan's per column?  
2 df_shop.isna().sum()
```

```
Out[23]: purchase_id      0  
lastname            0  
purchase_date       0  
stars              0  
price             22  
favorite_flower    178  
dtype: int64
```

```
In [24]: 1 # How many total nan's?  
2 df_shop.isna().sum().sum()
```

```
Out[24]: 200
```

Missing Data: Drop Rows

Missing Data: Drop Rows

```
In [25]: 1 df_shop.shape
```

```
Out[25]: (999, 6)
```

Missing Data: Drop Rows

```
In [25]: 1 df_shop.shape
```

```
Out[25]: (999, 6)
```

```
In [26]: 1 # drop rows with nan in any column
2 df_shop.dropna().shape
```

```
Out[26]: (801, 6)
```

Missing Data: Drop Rows

```
In [25]: 1 df_shop.shape
```

```
Out[25]: (999, 6)
```

```
In [26]: 1 # drop rows with nan in any column
2 df_shop.dropna().shape
```

```
Out[26]: (801, 6)
```

```
In [27]: 1 # drop only rows with nan in price using subset
2 df_shop.dropna(subset=['price']).shape
```

```
Out[27]: (977, 6)
```

Missing Data: Drop Rows

```
In [25]: 1 df_shop.shape
```

```
Out[25]: (999, 6)
```

```
In [26]: 1 # drop rows with nan in any column
2 df_shop.dropna().shape
```

```
Out[26]: (801, 6)
```

```
In [27]: 1 # drop only rows with nan in price using subset
2 df_shop.dropna(subset=['price']).shape
```

```
Out[27]: (977, 6)
```

```
In [28]: 1 # drop only rows with nans in all columns (a row of all nans)
2 df_shop.dropna(how='all').shape
```

```
Out[28]: (999, 6)
```

Missing Data: Drop Rows Cont.

Missing Data: Drop Rows Cont.

In [29]:

```
1 # save a new dataframe with dropped rows
2 df_shop = df_shop_raw.dropna().copy()
3 df_shop.shape
```

Out[29]: (803, 6)

Missing Data: Drop Rows Cont.

In [29]:

```
1 # save a new dataframe with dropped rows
2 df_shop = df_shop_raw.dropna().copy()
3 df_shop.shape
```

Out[29]: (803, 6)

In [30]:

```
1 # drop rows in current dataframe with inplace
2 df_shop = df_shop_raw.copy()
3
4 df_shop.dropna(inplace=True)
5 df_shop.shape
```

Out[30]: (803, 6)

Missing Data: Drop Rows

- Pros:
 - easy to do
 - simple to understand
- Cons:
 - potentially large data loss

Missing Data: Fill with Constant

- Use `.fillna()`
- Common filler: 0, -1

```
In [31]: 1 df_shop = df_shop_raw.drop_duplicates().copy() # make a new copy of the data
```

Missing Data: Fill with Constant

- Use `.fillna()`
- Common filler: 0, -1

```
In [31]: 1 df_shop = df_shop_raw.drop_duplicates().copy() # make a new copy of the data
```

```
In [32]: 1 df_shop.price[20:22]
```

```
Out[32]: 20      NaN
21    10.53
Name: price, dtype: float64
```

Missing Data: Fill with Constant

- Use `.fillna()`
- Common filler: 0, -1

```
In [31]: 1 df_shop = df_shop_raw.drop_duplicates().copy() # make a new copy of the data
```

```
In [32]: 1 df_shop.price[20:22]
```

```
Out[32]: 20      NaN  
21    10.53  
Name: price, dtype: float64
```

```
In [33]: 1 df_shop.price[20:22].fillna(0)
```

```
Out[33]: 20    0.00  
21    10.53  
Name: price, dtype: float64
```

Missing Data: Fill with Constant

- Use `.fillna()`
- Common filler: 0, -1

```
In [31]: 1 df_shop = df_shop_raw.drop_duplicates().copy() # make a new copy of the data
```

```
In [32]: 1 df_shop.price[20:22]
```

```
Out[32]: 20      NaN
          21    10.53
Name: price, dtype: float64
```

```
In [33]: 1 df_shop.price[20:22].fillna(0)
```

```
Out[33]: 20    0.00
          21  10.53
Name: price, dtype: float64
```

```
In [34]: 1 print(df_shop.price.mean().round(2))
2 print(df_shop.price.fillna(0).mean().round(2))
```

```
23.4
22.89
```

Missing Data: Fill with Constant

Pros:

- easy to do
- simple to understand

Cons:

- values may not be realistic

Missing Data: Impute

- Impute: fill with value inferred from existing values in that column
- Use `.fillna()` or sklearn methods
- Common filler values:
 - mean
 - median
 - "most frequent" aka mode

Missing Data: Impute

Missing Data: Impute

```
In [35]: 1 print(df_shop.price.mean().round(2))
2 print(df_shop.price.fillna(df_shop.price.mean()).mean().round(2))
```

```
23.4
23.4
```

Missing Data: Impute

In [35]:

```
1 print(df_shop.price.mean().round(2))
2 print(df_shop.price.fillna(df_shop.price.mean()).mean().round(2))
```

```
23.4
23.4
```

In [36]:

```
1 # make a copy to keep our original df
2 df_shop_impute = df_shop.copy()
```

Missing Data: Impute

```
In [35]: 1 print(df_shop.price.mean().round(2))
2 print(df_shop.price.fillna(df_shop.price.mean()).mean().round(2))
```

```
23.4
23.4
```

```
In [36]: 1 # make a copy to keep our original df
2 df_shop_impute = df_shop.copy()
```

```
In [37]: 1 # fill missing price with mean of price
2 df_shop_impute['price'] = df_shop.price.fillna(df_shop.price.mean())
```

Missing Data: Impute

```
In [35]: 1 print(df_shop.price.mean().round(2))
2 print(df_shop.price.fillna(df_shop.price.mean()).mean().round(2))
```

```
23.4
23.4
```

```
In [36]: 1 # make a copy to keep our original df
2 df_shop_impute = df_shop.copy()
```

```
In [37]: 1 # fill missing price with mean of price
2 df_shop_impute['price'] = df_shop.price.fillna(df_shop.price.mean())
```

```
In [38]: 1 # check to make sure all nulls filled
2 assert df_shop_impute.price.isna().sum() == 0
3 assert df_shop_impute.price.notna().all()
4
5 # also, that our mean hasn't changed
6 assert df_shop.price.mean() == df_shop_impute.price.mean()
```

```
In [112]: 1 assert df_shop.price.std() > df_shop_impute.price.std()
```

Missing Data: Impute

In [35]:

```
1 print(df_shop.price.mean().round(2))
2 print(df_shop.price.fillna(df_shop.price.mean()).mean().round(2))
```

```
23.4
23.4
```

In [36]:

```
1 # make a copy to keep our original df
2 df_shop_impute = df_shop.copy()
```

In [37]:

```
1 # fill missing price with mean of price
2 df_shop_impute['price'] = df_shop.price.fillna(df_shop.price.mean())
```

In [38]:

```
1 # check to make sure all nulls filled
2 assert df_shop_impute.price.isna().sum() == 0
3 assert df_shop_impute.price.notna().all()
4
5 # also, that our mean hasn't changed
6 assert df_shop.price.mean() == df_shop_impute.price.mean()
```

In [112]:

```
1 assert df_shop.price.std() > df_shop_impute.price.std()
```

In [39]:

```
1 # inplace works here as well
2 df_shop_impute.price.fillna(df_shop_impute.price.mean(), inplace=True)
```

Missing Data: Impute Cont.

if data is categorical?

Missing Data: Impute Cont.

if data is categorical?

```
In [40]: 1 df_shop.favorite_flower.mode() # may be more than 1!  
Out[40]: 0    lilac  
Name: favorite_flower, dtype: object
```

Missing Data: Impute Cont.

if data is categorical?

```
In [40]: 1 df_shop.favorite_flower.mode() # may be more than 1!
```

```
Out[40]: 0    lilac
Name: favorite_flower, dtype: object
```

```
In [41]: 1 # Note that we have to index into the DataFrame returned by mode to get a value
2 df_shop_impute.favorite_flower.fillna(df_shop_impute.favorite_flower.mode().iloc[0], inplace=True)
3
4 df_shop_impute.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1000 entries, 0 to 999
Data columns (total 6 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   purchase_id      1000 non-null   int64  
 1   lastname         1000 non-null   object  
 2   purchase_date    1000 non-null   datetime64[ns]
 3   stars            1000 non-null   int64  
 4   price            1000 non-null   float64 
 5   favorite_flower  1000 non-null   object  
dtypes: datetime64[ns](1), float64(1), int64(2), object(2)
memory usage: 87.0+ KB
```

Missing Data: Impute Cont. Using SimpleImputer

Missing Data: Impute Cont. Using SimpleImputer

```
In [42]: 1 df_shop[['price', 'stars']].loc[20:22]
```

Out[42]:

	price	stars
20	NaN	3
21	10.53	2
22	19.77	1

Missing Data: Impute Cont. Using SimpleImputer

```
In [42]: 1 df_shop[['price', 'stars']].loc[20:22]
```

Out[42]:

	price	stars
20	NaN	3
21	10.53	2
22	19.77	1

```
In [43]: 1 from sklearn.impute import SimpleImputer
```

```
2
```

```
3 imp = SimpleImputer(strategy='mean').fit(df_shop[['price', 'stars']])
```

```
4 print(f'fill values = {imp.statistics_.round(2)}')
```

```
5 imp.transform(df_shop.loc[20:22,['price', 'stars']]).round(2)
```

fill values = [23.4 3.6]

```
Out[43]: array([[23.4 , 3.  ],
   [10.53, 2.  ],
   [19.77, 1.  ]])
```

Missing Data: Impute Cont. Using SimpleImputer

```
In [42]: 1 df_shop[['price', 'stars']].loc[20:22]
```

Out[42]:

	price	stars
20	NaN	3
21	10.53	2
22	19.77	1

```
In [43]: 1 from sklearn.impute import SimpleImputer
```

```
2
```

```
3 imp = SimpleImputer(strategy='mean').fit(df_shop[['price', 'stars']])
```

```
4 print(f'fill values = {imp.statistics_.round(2)}')
```

```
5 imp.transform(df_shop.loc[20:22,['price', 'stars']]).round(2)
```

```
fill values = [23.4 3.6]
```

```
Out[43]: array([[23.4 , 3.  ],
   [10.53, 2.  ],
   [19.77, 1.  ]])
```

```
In [44]: 1 df_shop.favorite_flower[:3]
```

```
Out[44]: 0      iris
1      NaN
2  carnation
Name: favorite_flower, dtype: object
```

Missing Data: Impute Cont. Using SimpleImputer

```
In [42]: 1 df_shop[['price', 'stars']].loc[20:22]
```

Out[42]:

	price	stars
20	NaN	3
21	10.53	2
22	19.77	1

```
In [43]: 1 from sklearn.impute import SimpleImputer
```

```
2
```

```
3 imp = SimpleImputer(strategy='mean').fit(df_shop[['price', 'stars']])
```

```
4 print(f'fill values = {imp.statistics_.round(2)}')
```

```
5 imp.transform(df_shop.loc[20:22,['price', 'stars']]).round(2)
```

fill values = [23.4 3.6]

```
Out[43]: array([[23.4 , 3.  ],
   [10.53, 2.  ],
   [19.77, 1.  ]])
```

```
In [44]: 1 df_shop.favorite_flower[:3]
```

```
Out[44]: 0      iris
1      NaN
2  carnation
Name: favorite_flower, dtype: object
```

```
In [45]: 1 imp = SimpleImputer(strategy='most_frequent').fit(df_shop[['favorite_flower']])
2 imp.transform(df_shop.loc[:2,['favorite_flower']])
```

```
Out[45]: array([['iris'],
   ['lilac'],
   ['carnation']], dtype=object)
```

Missing Data: Impute

- Pros:
 - easy to do
 - simple to understand
- Cons:
 - may missing feature interactions

Missing Data: Infer

- Predict values of missing features using a model
- Ex: Can we predict price based on any of the other features?
- Additional feature engineering may be needed prior to this

Missing Data: Infer

- Predict values of missing features using a model
- Ex: Can we predict price based on any of the other features?
- Additional feature engineering may be needed prior to this

```
In [46]: 1 from sklearn.linear_model import LinearRegression
2
3 df_shop_infer = df_shop.copy()
4
5 not_missing = df_shop_infer.price.notna()
6 missing     = df_shop_infer.price.isna()
7
8 lr = LinearRegression().fit(df_shop_infer.loc[not_missing,['stars']],
9                             df_shop_infer.loc[not_missing,'price'])
10
11 df_shop_infer.loc[missing,'price'] = lr.predict(df_shop_infer.loc[missing,['stars']])
```

Missing Data: Infer

- Pros:
 - better estimate (based on other data)
- Cons:
 - have to train another model
 - colinear features!

Missing Data: Adjacent Data

- Use `.fillna()` with method:
 - `ffill`: propagate last valid observation forward to next valid
 - `bfill`: use next valid observation to fill gap backwards
- Use when there is reason to believe data not i.i.d. (eg: **timeseries**)

Missing Data: Adjacent Data

- Use `.fillna()` with method:
 - `ffill`: propagate last valid observation forward to next valid
 - `bfill`: use next valid observation to fill gap backwards
- Use when there is reason to believe data not i.i.d. (eg: **timeseries**)

```
In [47]: 1 df_shop.price.loc[19:21]
```

```
Out[47]: 19    20.45
20      NaN
21    10.53
Name: price, dtype: float64
```

Missing Data: Adjacent Data

- Use `.fillna()` with method:
 - `ffill`: propagate last valid observation forward to next valid
 - `bfill`: use next valid observation to fill gap backwards
- Use when there is reason to believe data not i.i.d. (eg: **timeseries**)

```
In [47]: 1 df_shop.price.loc[19:21]
```

```
Out[47]: 19    20.45
20      NaN
21    10.53
Name: price, dtype: float64
```

```
In [48]: 1 df_shop.price.fillna(method='ffill').loc[19:21]
```

```
Out[48]: 19    20.45
20    20.45
21    10.53
Name: price, dtype: float64
```

Missing Data: Add a Dummy Column First

- Data may be missing for a reason!
- Capture "missing" before filling

Missing Data: Add a Dummy Column First

- Data may be missing for a reason!
- Capture "missing" before filling

```
In [49]: 1 df_shop = df_shop_raw.drop_duplicates().copy()  
2  
3 # storing a column of 1:missing, 0:not-missing  
4 df_shop['price_missing'] = df_shop.price.isna().astype(int)  
5  
6 # can now fill missing values  
7 df_shop['price'] = df_shop.price.fillna(df_shop.price.mean())
```

Missing Data: Add a Dummy Column First

- Data may be missing for a reason!
- Capture "missing" before filling

```
In [49]: 1 df_shop = df_shop_raw.drop_duplicates().copy()
2
3 # storing a column of 1:missing, 0:not-missing
4 df_shop['price_missing'] = df_shop.price.isna().astype(int)
5
6 # can now fill missing values
7 df_shop['price'] = df_shop.price.fillna(df_shop.price.mean())
```

```
In [50]: 1 # finding where data was missing
2 np.where(df_shop.price_missing == 1)
```

```
Out[50]: (array([ 20,  41,  54,  63, 145, 186, 194, 203, 212, 360, 367, 382, 429,
   469, 522, 570, 595, 726, 792, 821, 974, 978]),)
```

Missing Data: Add a Dummy Column First

- Data may be missing for a reason!
- Capture "missing" before filling

```
In [49]: 1 df_shop = df_shop_raw.drop_duplicates().copy()
2
3 # storing a column of 1:missing, 0:not-missing
4 df_shop['price_missing'] = df_shop.price.isna().astype(int)
5
6 # can now fill missing values
7 df_shop['price'] = df_shop.price.fillna(df_shop.price.mean())
```

```
In [50]: 1 # finding where data was missing
2 np.where(df_shop.price_missing == 1)
```

```
Out[50]: (array([ 20,  41,  54,  63, 145, 186, 194, 203, 212, 360, 367, 382, 429,
 469, 522, 570, 595, 726, 792, 821, 974, 978]),)
```

```
In [51]: 1 df_shop[['price', 'price_missing']].iloc[20:23]
```

```
Out[51]:
  price  price_missing
20  23.403384    1
21  10.530000    0
22  19.770000    0
```

Rescaling

- Often need features to be in the same scale
- Methods of rescaling
 - Standardization (z-score)
 - Min-Max rescaling
 - others...

Rescaling

- Often need features to be in the same scale
- Methods of rescaling
 - Standardization (z-score)
 - Min-Max rescaling
 - others...

```
In [52]: 1 # load taxi data
2 df_taxi_raw = pd.read_csv('../data/yellowcab_tripdata_2017-01_subset10000rows.csv',
3                         parse_dates=['tpep_pickup_datetime', 'tpep_dropoff_datetime'])
4 # create trip_duration
5 df_taxi_raw['trip_duration'] = (df_taxi_raw.tpep_dropoff_datetime - df_taxi_raw.tpep_pickup_datetime).dt.seconds
6
7 # select subset
8 df_taxi_raw = df_taxi_raw[df_taxi_raw.trip_duration.lt(3600) &
9                         df_taxi_raw.tip_amount.between(0,10,inclusive='neither')]
10
11 df_taxi = df_taxi_raw.copy()
```

Rescaling

- Often need features to be in the same scale
- Methods of rescaling
 - Standardization (z-score)
 - Min-Max rescaling
 - others...

```
In [52]: 1 # load taxi data
2 df_taxi_raw = pd.read_csv('../data/yellowcab_tripdata_2017-01_subset10000rows.csv',
3                           parse_dates=['tpep_pickup_datetime', 'tpep_dropoff_datetime'])
4 # create trip_duration
5 df_taxi_raw['trip_duration'] = (df_taxi_raw.tpep_dropoff_datetime - df_taxi_raw.tpep_pickup_datetime).dt.seconds
6
7 # select subset
8 df_taxi_raw = df_taxi_raw[df_taxi_raw.trip_duration.lt(3600) &
9                           df_taxi_raw.tip_amount.between(0,10,inclusive='neither')]
10
11 df_taxi = df_taxi_raw.copy()
```

```
In [53]: 1 df_taxi[['trip_duration', 'tip_amount']].agg(['mean', 'std', 'min', 'max'], axis=0).round(2)
```

Out[53]:

	trip_duration	tip_amount
mean	765.03	2.41
std	496.83	1.55
min	2.00	0.01

Rescaling: Standardization

- rescale to 0 mean, standard deviation of 1
 - $X_{\text{scaled}} = (X - X.\text{mean}) / X.\text{std}$ ()

Rescaling: Standardization

- rescale to 0 mean, standard deviation of 1
 - $X_{\text{scaled}} = (X - X.\text{mean}) / X.\text{std}$ ()

```
In [54]: 1 from sklearn.preprocessing import StandardScaler  
2  
3 # instantiate  
4 ss = StandardScaler(with_mean=True, with_std=True) # default is center and scale  
5  
6 # fit to the data  
7 ss.fit(df_taxi[['trip_duration','tip_amount']])  
8  
9 # transform the data  
10 X_ss = ss.transform(df_taxi[['trip_duration','tip_amount']])  
11 X_ss[:2].round(2)
```

```
Out[54]: array([[-0.5 , -0.48],  
                 [-0.17, -0.91]])
```

Rescaling: Standardization

- rescale to 0 mean, standard deviation of 1
 - $X_{\text{scaled}} = (X - X.\text{mean}) / X.\text{std}()$

```
In [54]: 1 from sklearn.preprocessing import StandardScaler  
2  
3 # instantiate  
4 ss = StandardScaler(with_mean=True, with_std=True) # default is center and scale  
5  
6 # fit to the data  
7 ss.fit(df_taxi[['trip_duration','tip_amount']])  
8  
9 # transform the data  
10 X_ss = ss.transform(df_taxi[['trip_duration','tip_amount']])  
11 X_ss[:2].round(2)
```

```
Out[54]: array([[-0.5 , -0.48],  
                 [-0.17, -0.91]])
```

```
In [55]: 1 df_taxi_ss = pd.DataFrame(X_ss,columns=['trip_duration_scaled','tip_amount_scaled'])  
2 df_taxi_ss.agg(['mean','std','min','max'],axis=0).round(2)
```

```
Out[55]:
```

	trip_duration_scaled	tip_amount_scaled
mean	0.00	-0.00
std	1.00	1.00
min	-1.54	-1.54
max	5.62	4.88

Rescaling: Min-Max

- rescale values between 0 and 1
- $X_{\text{scaled}} = (X - X.\min()) / (X.\max() - X.\min())$
- removes negative values

Rescaling: Min-Max

- rescale values between 0 and 1
- $X_{\text{scaled}} = (X - X.\min()) / (X.\max() - X.\min())$
- removes negative values

```
In [56]: 1 from sklearn.preprocessing import MinMaxScaler  
2  
3 X_mms = MinMaxScaler(feature_range=(0,1) # default is to rescale between 0 and 1  
4         .fit_transform(df_taxi[['trip_duration','tip_amount']])  
5  
6 df_taxi_mms = pd.DataFrame(X_mms,columns=['trip_duration_scaled','tip_amount_scaled'])  
7 df_taxi_mms.agg(['mean','std','min','max']).round(2)
```

Out[56]:

	trip_duration_scaled	tip_amount_scaled
mean	0.21	0.24
std	0.14	0.16
min	0.00	0.00
max	1.00	1.00




Dealing with Skew

- Many models expect "normal", symmetric data (ex: linear models)
- Highly skewed: tail has larger effect on model (outliers?)
- Transform with `log` or `sqrt`

Dealing with Skew


- Many models expect "normal", symmetric data (ex: linear models)
- Highly skewed: tail has larger effect on model (outliers?)
- Transform with `log` or `sqrt`

```
In [57]: 1 fig,ax = plt.subplots(1,3,figsize=(16,4))
2 sns.histplot(x=df_taxi.total_amount, ax=ax[0]);
3 sns.histplot(x=df_taxi.total_amount.apply(np.sqrt), ax=ax[1]); ax[1].set_xlabel('sqrt transform');
4 sns.histplot(x=df_taxi.total_amount.apply(np.log),   ax=ax[2]); ax[2].set_xlabel('log transform');
```



Outliers

- Similar to missing data:
 - human data entry error
 - instrument measurement errors
 - data processing errors
 - natural deviations



Outliers



- Why worry about them?
 - can give misleading results
 - can indicate issues in data/measurement
- Detecting Outliers
 - understand your data!
 - visualizations
 - $1.5 * \text{IQR}$
 - z-scores
 - etc..

Detecting Outliers

Detecting Outliers

In [58]:


```
1 np.random.seed(123)
2 data_rand = np.concatenate([np.random.normal(0,2,200),np.random.normal(20,.2,20)])
3 df_rand = pd.DataFrame({'measure':data_rand})
4
5 fig,ax = plt.subplots(1,2, figsize=(14,4))
6 sns.histplot(x=df_rand.measure,ax=ax[0]);sns.boxplot(x=df_rand.measure,ax=ax[1]);
```



Detecting Outliers

In [58]:

```
1 np.random.seed(123)
2 data_rand = np.concatenate([np.random.normal(0, 2, 200), np.random.normal(20, .2, 20)])
3 df_rand = pd.DataFrame({'measure':data_rand})
4
5 fig,ax = plt.subplots(1,2, figsize=(14,4))
6 sns.histplot(x=df_rand.measure,ax=ax[0]);sns.boxplot(x=df_rand.measure,ax=ax[1]);
```



In [59]:


```
1 # Calculating IQR
2 p25,p75 = df_rand.measure.quantile(.25,.75)
3 iqr = p75 - p25
4 round(iqr,2)
```

Out[59]: 3.3

Detecting Outliers

In [58]:

```
1 np.random.seed(123)
2 data_rand = np.concatenate([np.random.normal(0, 2, 200), np.random.normal(20, .2, 20)])
3 df_rand = pd.DataFrame({'measure':data_rand})
4
5 fig,ax = plt.subplots(1,2, figsize=(14,4))
6 sns.histplot(x=df_rand.measure,ax=ax[0]);sns.boxplot(x=df_rand.measure,ax=ax[1]);
```



In [59]:

```
1 # Calculating IQR
2 p25,p75 = df_rand.measure.quantile([.25,.75])
3 iqr = p75 - p25
4 round(iqr,2)
```

Out[59]: 3.3

In [60]:

```
1 # Finding outliers with IQR (first two examples found)
2 df_rand.measure[(df_rand.measure > p75+(1.5*iqr)) | (df_rand.measure < p25-(1.5*iqr))].sort_values()#.head(2).round(2)
```

Detecting Outliers with z-score




Detecting Outliers with z-score

In [61]:

```
1 # zscore
2 df_rand['measure_zscore'] = (df_rand.measure - df_rand.measure.mean()) / df_rand.measure.std()
3
4 fig, ax = plt.subplots(1,3,figsize=(16,4))
5 sns.histplot(x=df_rand.measure,ax=ax[0]);
6 sns.histplot(x=df_rand.measure_zscore, ax=ax[1]);
7
8 keep_idx = np.abs(df_rand.measure_zscore) < 2
9 sns.histplot(x=df_rand[keep_idx].measure_zscore, ax=ax[2]);
10
11 # sample of points getting dropped
12 df_rand[np.abs(df_rand.measure_zscore) >= 2].sort_values(by='measure').head(3).round(2)
```

Out[61]:

	measure	measure_zscore
213	19.72	2.93
207	19.82	2.94
218	19.85	2.95



Other Outlier Detection Methods

- Many more parametric and non-parametric methods
 - Standardized Residuals
 - DBScan
 - EllipticEnvelope
 - IsolationForest
 - other Anomaly Detection techniques
 - See [sklearn docs on Outlier Detection](#) for more details

Dealing with Outliers

- How to deal with outliers?
 - drop data
 - treat as missing
 - encode with dummy variable first

Putting It All Together: Different Styles

Putting It All Together: Different Styles

```
In [62]: 1 df_shop1 = pd.read_csv('../data/flowershop_data_with_dups_week8.csv')
2 df_shop1 = df_shop1.drop_duplicates()
3 df_shop1['purchase_date'] = pd.to_datetime(df_shop1.purchase_date)
4 df_shop1['price_missing'] = df_shop1.price.isna().astype(int)
5 df_shop1['price'] = df_shop1.price.fillna(df_shop1.price.mean())
6 df_shop1['price_scaled'] = StandardScaler().fit_transform(df_shop1[['price']])
7 df_shop1['favorite_flower_missing'] = df_shop1.favorite_flower.isna().astype(int)
8 df_shop1['favorite_flower'] = SimpleImputer(strategy='most_frequent').fit(df_shop1[['favorite_flower']])
```

Putting It All Together: Different Styles

```
In [62]: 1 df_shop1 = pd.read_csv('../data/flowershop_data_with_dups_week8.csv')
2 df_shop1 = df_shop1.drop_duplicates()
3 df_shop1['purchase_date'] = pd.to_datetime(df_shop1.purchase_date)
4 df_shop1['price_missing'] = df_shop1.price.isna().astype(int)
5 df_shop1['price'] = df_shop1.price.fillna(df_shop1.price.mean())
6 df_shop1['price_scaled'] = StandardScaler().fit_transform(df_shop1[['price']])
7 df_shop1['favorite_flower_missing'] = df_shop1.favorite_flower.isna().astype(int)
8 df_shop1['favorite_flower'] = SimpleImputer(strategy='most_frequent').fit(df_shop1[['favorite_flower']])
```

```
In [63]: 1 df_shop2 = (
2     pd.read_csv('../data/flowershop_data_with_dups_week8.csv')
3     .drop_duplicates()
4     .assign(
5         purchase_date = lambda df_ : pd.to_datetime(df_.purchase_date),
6         price_missing = lambda df_ : df_.price.isna().astype(int),
7         price = lambda df_ : df_.price.fillna(df_.price.mean()),
8         price_scaled = lambda df_ : StandardScaler().fit_transform(df_[['price']]),
9         favorite_flower_missing = lambda df_ : df_.favorite_flower.isna().astype(int),
10        favorite_flower = lambda df_ : (SimpleImputer(strategy='most_frequent')
11                                         .fit_transform(df_shop1[['favorite_flower']]))
12    )
13 )
14 )
```

Putting It All Together: Different Styles

```
In [62]: 1 df_shop1 = pd.read_csv('../data/flowershop_data_with_dups_week8.csv')
2 df_shop1 = df_shop1.drop_duplicates()
3 df_shop1['purchase_date'] = pd.to_datetime(df_shop1.purchase_date)
4 df_shop1['price_missing'] = df_shop1.price.isna().astype(int)
5 df_shop1['price'] = df_shop1.price.fillna(df_shop1.price.mean())
6 df_shop1['price_scaled'] = StandardScaler().fit_transform(df_shop1[['price']])
7 df_shop1['favorite_flower_missing'] = df_shop1.favorite_flower.isna().astype(int)
8 df_shop1['favorite_flower'] = SimpleImputer(strategy='most_frequent').fit(df_shop1[['favorite_flower']])
```

```
In [63]: 1 df_shop2 = (
2     pd.read_csv('../data/flowershop_data_with_dups_week8.csv')
3     .drop_duplicates()
4     .assign(
5         purchase_date = lambda df_ : pd.to_datetime(df_.purchase_date),
6         price_missing = lambda df_ : df_.price.isna().astype(int),
7         price = lambda df_ : df_.price.fillna(df_.price.mean()),
8         price_scaled = lambda df_ : StandardScaler().fit_transform(df_[['price']]),
9         favorite_flower_missing = lambda df_ : df_.favorite_flower.isna().astype(int),
10        favorite_flower = lambda df_ : (SimpleImputer(strategy='most_frequent')
11                                         .fit_transform(df_shop1[['favorite_flower']]))
12    )
13 )
14 )
```

```
In [64]: 1 pd.testing.assert_frame_equal(df_shop1,df_shop2) # throws an exception when data frames are not the same
```

Data Cleaning Review

- duplicate data
- missing data
- rescaling
- dealing with skew
- outlier detection

Feature Engineering

- Binning
- One-Hot encoding
- Derived Features

Binning

- Transform continuous features to categorical
- Use:
 - `pd.cut`
 - `sklearn.preprocessing.KBinsDiscretizer` (combined binning and one-hot-encoding)

Binning

- Transform continuous features to categorical
- Use:
 - pd.cut
 - sklearn.preprocessing.KBinsDiscretizer (combined binning and one-hot-encoding)

```
In [65]: 1 trip_duration_bins = [df_taxi.trip_duration.min(),
 2                         df_taxi.trip_duration.median(),
 3                         df_taxi.trip_duration.quantile(0.75),
 4                         df_taxi.trip_duration.max(),]
```

Binning

- Transform continuous features to categorical
- Use:
 - `pd.cut`
 - `sklearn.preprocessing.KBinsDiscretizer` (combined binning and one-hot-encoding)

```
In [65]: 1 trip_duration_bins = [df_taxi.trip_duration.min(),
2                           df_taxi.trip_duration.median(),
3                           df_taxi.trip_duration.quantile(0.75),
4                           df_taxi.trip_duration.max(),]
```

```
In [66]: 1 df_taxi_bin = df_taxi_raw.copy()
2 df_taxi_bin['trip_duration_binned'] = pd.cut(df_taxi_bin.trip_duration,
3                                              bins=trip_duration_bins,           # can pass bin edges or number of bins
4                                              labels=['short','medium','long'],
5                                              right=True,                      # all bins right-inclusive
6                                              include_lowest=True)            # first interval left-inclusive
7
8 df_taxi_bin[['trip_duration','trip_duration_binned']].iloc[:10]
9
```

Out[66]:

	trip_duration	trip_duration_binned
1	516	short
2	683	medium
7	834	medium
8	298	short

One-Hot Encoding

- Encode categorical features for models that can't handle categorical (eg. Linear)
- One column per category, '1' in only one column per row
- Use `pd.get_dummies()` or `sklearn.preprocessing.OneHotEncoder`

One-Hot Encoding

- Encode categorical features for models that can't handle categorical (eg. Linear)
- One column per category, '1' in only one column per row
- Use `pd.get_dummies()` or `sklearn.preprocessing.OneHotEncoder`

```
In [67]: 1 pd.get_dummies(df_taxi_bin.trip_duration_binned, prefix='trip_duration').iloc[:2]
```

Out[67]:

	trip_duration_short	trip_duration_medium	trip_duration_long
1	1	0	0
2	0	1	0

One-Hot Encoding

- Encode categorical features for models that can't handle categorical (eg. Linear)
- One column per category, '1' in only one column per row
- Use `pd.get_dummies()` or `sklearn.preprocessing.OneHotEncoder`

```
In [67]: 1 pd.get_dummies(df_taxi_bin.trip_duration_binned, prefix='trip_duration').iloc[:2]
```

Out[67]:

	trip_duration_short	trip_duration_medium	trip_duration_long
1	1	0	0
2	0	1	0

```
In [68]: 1 # to add back to dataframe, use join (will discuss .join() next time)
```

```
2 df_taxi_bin.join(pd.get_dummies(df_taxi_bin.trip_duration_binned, prefix='trip_duration')).iloc[:2,-6:] # not saved
```

Out[68]:

	total_amount	trip_duration	trip_duration_binned	trip_duration_short	trip_duration_medium	trip_duration_long
1	9.96	516	short	1	0	0
2	10.30	683	medium	0	1	0

One-Hot Encoding

- Encode categorical features for models that can't handle categorical (eg. Linear)
- One column per category, '1' in only one column per row
- Use `pd.get_dummies()` or `sklearn.preprocessing.OneHotEncoder`

```
In [67]: 1 pd.get_dummies(df_taxi_bin.trip_duration_binned, prefix='trip_duration').iloc[:2]
```

Out[67]:

	trip_duration_short	trip_duration_medium	trip_duration_long
1	1	0	0
2	0	1	0

```
In [68]: 1 # to add back to dataframe, use join (will discuss .join() next time)
```

```
2 df_taxi_bin.join(pd.get_dummies(df_taxi_bin.trip_duration_binned, prefix='trip_duration')).iloc[:2,-6:] # not saved
```

Out[68]:

	total_amount	trip_duration	trip_duration_binned	trip_duration_short	trip_duration_medium	trip_duration_long
1	9.96	516	short	1	0	0
2	10.30	683	medium	0	1	0

```
In [69]: 1 # or let pandas determine which columns to one-hot
```

```
2 pd.get_dummies(df_taxi_bin).iloc[:2,-6:] # not being saved
```

Out[69]:

	trip_duration	store_and_fwd_flag_N	store_and_fwd_flag_Y	trip_duration_binned_short	trip_duration_binned_medium	trip_duration_binned_long
1	516	1	0	1	0	0
2	683	1	0	0	1	0

One-Hot Encoding with sklearn

One-Hot Encoding with sklearn

```
In [70]: 1 from sklearn.preprocessing import OneHotEncoder  
2  
3 ohe = OneHotEncoder(categories=[[ 'short', 'medium', 'long']], # or leave as 'auto'  
4           sparse=True,  
5           handle_unknown='ignore')                         # will raise error otherwise  
6  
7 ohe.fit(df_taxi_bin[['trip_duration_binned']])  
8 ohe.categories_
```

```
Out[70]: [array(['short', 'medium', 'long'], dtype=object)]
```

One-Hot Encoding with sklearn

```
In [70]: 1 from sklearn.preprocessing import OneHotEncoder  
2  
3 ohe = OneHotEncoder(categories=[[ 'short', 'medium', 'long']], # or leave as 'auto'  
4           sparse=True,  
5           handle_unknown='ignore') # will raise error otherwise  
6  
7 ohe.fit(df_taxi_bin[['trip_duration_binned']])  
8 ohe.categories_
```

```
Out[70]: [array(['short', 'medium', 'long'], dtype=object)]
```

```
In [71]: 1 ohe.transform(df_taxi_bin[['trip_duration_binned']])[:3] # returns a sparse matrix!
```

```
Out[71]: <3x3 sparse matrix of type '<class 'numpy.float64'>'  
with 3 stored elements in Compressed Sparse Row format>
```

One-Hot Encoding with sklearn

```
In [70]: 1 from sklearn.preprocessing import OneHotEncoder  
2  
3 ohe = OneHotEncoder(categories=[[ 'short', 'medium', 'long']], # or leave as 'auto'  
4           sparse=True,  
5           handle_unknown='ignore') # will raise error otherwise  
6  
7 ohe.fit(df_taxi_bin[['trip_duration_binned']])  
8 ohe.categories_
```

```
Out[70]: [array(['short', 'medium', 'long'], dtype=object)]
```

```
In [71]: 1 ohe.transform(df_taxi_bin[['trip_duration_binned']])[:3] # returns a sparse matrix!
```

```
Out[71]: <3x3 sparse matrix of type '<class 'numpy.float64'>'  
with 3 stored elements in Compressed Sparse Row format>
```

```
In [72]: 1 ohe.transform(df_taxi_bin[['trip_duration_binned']])[:3].todense() # use .todense() to convert sparse to dense
```

```
Out[72]: matrix([[1., 0., 0.],  
                 [0., 1., 0.],  
                 [0., 1., 0.]])
```

Bin and One-Hot Encode with sklearn

Bin and One-Hot Encode with sklearn

In [73]:

```
1 from sklearn.preprocessing import KBinsDiscretizer
2
3 # NOTE: We're not setting the bin edges explicitly
4 #       For control over bin edges, use Binarizer
5 kbd = KBinsDiscretizer(n_bins=3,
6                       encode="onehot",      # or onehot (sparse), ordinal
7                       strategy="quantile", # or uniform or kmeans (clustering)
8                       ).fit(df_taxi[['trip_duration']])
9 print(kbd.bin_edges_)
10 print(kbd.bin_edges_[0].astype(int))
```



```
[array([2.000e+00, 4.780e+02, 8.700e+02, 3.556e+03]])
[    2   478   870  3556]
```

Bin and One-Hot Encode with sklearn

In [73]:

```
1 from sklearn.preprocessing import KBinsDiscretizer
2
3 # NOTE: We're not setting the bin edges explicitly
4 #       For control over bin edges, use Binarizer
5 kbd = KBinsDiscretizer(n_bins=3,
6                       encode="onehot",      # or onehot (sparse), ordinal
7                       strategy="quantile", # or uniform or kmeans (clustering)
8                       ).fit(df_taxi[['trip_duration']])
9 print(kbd.bin_edges_)
10 print(kbd.bin_edges_[0].astype(int))
```

```
[array([2.000e+00, 4.780e+02, 8.700e+02, 3.556e+03])]
[    2   478   870  3556]
```

In [74]:

```
1 df_taxi[['trip_duration']].tail(3)
```

Out[74]:

	trip_duration
9994	905
9995	296
9997	2089

Bin and One-Hot Encode with sklearn

In [73]:

```
1 from sklearn.preprocessing import KBinsDiscretizer
2
3 # NOTE: We're not setting the bin edges explicitly
4 #       For control over bin edges, use Binarizer
5 kbd = KBinsDiscretizer(n_bins=3,
6                       encode="onehot",      # or onehot (sparse), ordinal
7                       strategy="quantile", # or uniform or kmeans (clustering)
8                       ).fit(df_taxi[['trip_duration']])
9 print(kbd.bin_edges_)
10 print(kbd.bin_edges_[0].astype(int))
```

```
[array([2.000e+00, 4.780e+02, 8.700e+02, 3.556e+03])]
[ 2 478 870 3556]
```

In [74]:

```
1 df_taxi[['trip_duration']].tail(3)
```

Out[74]:

	trip_duration
9994	905
9995	296
9997	2089

In [75]:

```
1 kbd.transform(df_taxi[['trip_duration']])[-3:]
```

Out[75]:

```
<3x3 sparse matrix of type '<class 'numpy.float64'>'  
with 3 stored elements in Compressed Sparse Row format>
```

Bin and One-Hot Encode with sklearn

In [73]:

```
1 from sklearn.preprocessing import KBinsDiscretizer
2
3 # NOTE: We're not setting the bin edges explicitly
4 #       For control over bin edges, use Binarizer
5 kbd = KBinsDiscretizer(n_bins=3,
6                       encode="onehot",      # or onehot (sparse), ordinal
7                       strategy="quantile", # or uniform or kmeans (clustering)
8                       ).fit(df_taxi[['trip_duration']])
9 print(kbd.bin_edges_)
10 print(kbd.bin_edges_[0].astype(int))
```

```
[array([2.000e+00, 4.780e+02, 8.700e+02, 3.556e+03])]
[ 2 478 870 3556]
```

In [74]:

```
1 df_taxi[['trip_duration']].tail(3)
```

Out[74]:

	trip_duration
9994	905
9995	296
9997	2089

In [75]:

```
1 kbd.transform(df_taxi[['trip_duration']])[-3:]
```

Out[75]:

```
<3x3 sparse matrix of type '<class 'numpy.float64'>'  
with 3 stored elements in Compressed Sparse Row format>
```

In [76]:

```
1 kbd.transform(df_taxi[['trip_duration']])[-5:].todense()
```

```
Out[76]: matrix([[0., 0., 1.],
                  [0., 0., 1.],
                  [0., 0., 1.],
                  [1., 0., 0.],
```

Dealing with Ordinal Variables

Dealing with Ordinal Variables

In [77]:

```
1 df_pml = pd.DataFrame([['green', 'M', 10.1, 'class2'],
2                         ['red', 'L', 13.5, 'class1'],
3                         ['blue', 'XL', 15.3, 'class2']],
4                         columns=['color', 'size', 'price', 'classlabel'])
5 df_pml
```

Out[77]:

	color	size	price	classlabel
0	green	M	10.1	class2
1	red	L	13.5	class1
2	blue	XL	15.3	class2

Dealing with Ordinal Variables

```
In [77]: 1 df_pml = pd.DataFrame([['green', 'M', 10.1, 'class2'],
2                           ['red', 'L', 13.5, 'class1'],
3                           ['blue', 'XL', 15.3, 'class2']],
4                           columns=['color', 'size', 'price', 'classlabel'])
5 df_pml
```

```
Out[77]:   color  size  price  classlabel
0  green    M   10.1  class2
1   red     L   13.5  class1
2  blue    XL   15.3  class2
```

```
In [78]: 1 # if we know the numerical difference between ordinal values
2 # eg XL = L+1 = M+2
3
4 size_mapping = {'XL':3,
5                  'L':2,
6                  'M':1}
7
8 df_pml_features = pd.DataFrame()
9
10 df_pml_features['size'] = df_pml['size'].map(size_mapping)
11 df_pml_features
```

```
Out[78]:  size
0   1
1   2
2   3
```

Dealing with Ordinal Variables Cont.

Dealing with Ordinal Variables Cont.

```
In [79]: 1 df_pml
```

```
Out[79]:
```

	color	size	price	classlabel
0	green	M	10.1	class2
1	red	L	13.5	class1
2	blue	XL	15.3	class2

Dealing with Ordinal Variables Cont.

In [79]:

```
1 df_pml
```

Out[79]:

	color	size	price	classlabel
0	green	M	10.1	class2
1	red	L	13.5	class1
2	blue	XL	15.3	class2

In [80]:

```
1 # if we don't know the numerical difference between ordinal values
2 # generate threshold features
3 df_pml_features = pd.DataFrame()
4 df_pml_features['x > M'] = df_pml['size'].apply(lambda x: 1 if x in ['L', 'XL'] else 0)
5 df_pml_features['x > L'] = df_pml['size'].apply(lambda x: 1 if x == 'XL' else 0)
6 df_pml_features
```

Out[80]:

	x > M	x > L
0	0	0
1	1	0
2	1	1

Derived Features

- Anything that is a transformation of our data
- This is where the money is!
- Examples:
 - "is a high demand pickup location"
 - "is a problem house sale"
 - "high-performing job candidate"

Polynomial Features

Polynomial Features

```
In [117]: 1 from sklearn.preprocessing import PolynomialFeatures  
2  
3 pf = PolynomialFeatures(degree=3,  
4                         include_bias=False)  
5 X_new = pf.fit_transform(df_taxi[['passenger_count','trip_duration']])  
6  
7 # new_columns = ['passenger_count','trip_duration','passenger_count^2','passenger_count*trip_duration','trip_duration^2']  
8 # pd.DataFrame(X_new[3:5],columns=new_columns)
```

```
In [118]: 1 X_new.shape
```

```
Out[118]: (6225, 9)
```

Python String Functions

Python String Functions

```
In [82]: 1 doc = "D.S. is good!"  
         2 doc
```

```
Out[82]: 'D.S. is good!'
```

Python String Functions

```
In [82]: 1 doc = "D.S. is good!"  
2 doc
```

```
Out[82]: 'D.S. is good!'
```

```
In [83]: 1 doc.lower(),doc.upper()      # change capitalization
```

```
Out[83]: ('d.s. is good!', 'D.S. IS GOOD!')
```

Python String Functions

```
In [82]: 1 doc = "D.S. is good!"  
2 doc
```

```
Out[82]: 'D.S. is good!'
```

```
In [83]: 1 doc.lower(),doc.upper()      # change capitalization
```

```
Out[83]: ('d.s. is good!', 'D.S. IS GOOD!')
```

```
In [84]: 1 doc.split() , doc.split('.')  # split a string into parts (default is whitespace)
```

```
Out[84]: ([ 'D.S.', 'is', 'good!' ], [ 'D', 'S', ' ' is good!] )
```

Python String Functions

```
In [82]: 1 doc = "D.S. is good!"  
2 doc
```

```
Out[82]: 'D.S. is good!'
```

```
In [83]: 1 doc.lower(),doc.upper()      # change capitalization
```

```
Out[83]: ('d.s. is good!', 'D.S. IS GOOD!')
```

```
In [84]: 1 doc.split() , doc.split('.')  # split a string into parts (default is whitespace)
```

```
Out[84]: ([ 'D.S.', 'is', 'good!' ], [ 'D', 'S', ' ' is good!] )
```

```
In [85]: 1 '|'.join(['ab','c','d'])      # join items in a list together
```

```
Out[85]: 'ab|c|d'
```

Python String Functions

```
In [82]: 1 doc = "D.S. is good!"  
2 doc
```

```
Out[82]: 'D.S. is good!'
```

```
In [83]: 1 doc.lower(),doc.upper()      # change capitalization
```

```
Out[83]: ('d.s. is good!', 'D.S. IS GOOD!')
```

```
In [84]: 1 doc.split() , doc.split('.')  # split a string into parts (default is whitespace)
```

```
Out[84]: (['D.S.', 'is', 'good!'], ['D', 'S', ' is good!'])
```

```
In [85]: 1 '|'.join(['ab','c','d'])      # join items in a list together
```

```
Out[85]: 'ab|c|d'
```

```
In [86]: 1 '|'.join(doc[:5])          # a string itself is treated like a list of characters
```

```
Out[86]: 'D|.|s|.| '
```

Python String Functions

```
In [82]: 1 doc = "D.S. is good!"  
2 doc
```

```
Out[82]: 'D.S. is good!'
```

```
In [83]: 1 doc.lower(),doc.upper()      # change capitalization
```

```
Out[83]: ('d.s. is good!', 'D.S. IS GOOD!')
```

```
In [84]: 1 doc.split() , doc.split('.')  # split a string into parts (default is whitespace)
```

```
Out[84]: (['D.S.', 'is', 'good!'], ['D', 'S', ' is good!'])
```

```
In [85]: 1 '|'.join(['ab','c','d'])      # join items in a list together
```

```
Out[85]: 'ab|c|d'
```

```
In [86]: 1 '|'.join(doc[:5])          # a string itself is treated like a list of characters
```

```
Out[86]: 'D|.|s|.| '
```

```
In [87]: 1 ' test '.strip()          # remove whitespace from the beginning and end of a string
```

```
Out[87]: 'test'
```

Python String Functions

```
In [82]: 1 doc = "D.S. is good!"  
2 doc
```

```
Out[82]: 'D.S. is good!'
```

```
In [83]: 1 doc.lower(),doc.upper()      # change capitalization
```

```
Out[83]: ('d.s. is good!', 'D.S. IS GOOD!')
```

```
In [84]: 1 doc.split() , doc.split('.')  # split a string into parts (default is whitespace)
```

```
Out[84]: (['D.S.', 'is', 'good!'], ['D', 'S', ' is good!'])
```

```
In [85]: 1 '|'.join(['ab','c','d'])      # join items in a list together
```

```
Out[85]: 'ab|c|d'
```

```
In [86]: 1 '|'.join(doc[:5])          # a string itself is treated like a list of characters
```

```
Out[86]: 'D|.|s|.| '
```

```
In [87]: 1 ' test '.strip()          # remove whitespace from the beginning and end of a string
```

```
Out[87]: 'test'
```

and more, see <https://docs.python.org/3.8/library/string.html>

String Functions in Pandas

String Functions in Pandas

```
In [88]: 1 df_shop.iloc[:2].loc[:, 'lastname']
```

```
Out[88]: 0    PERKINS
1    ROBINSON
Name: lastname, dtype: object
```

String Functions in Pandas

```
In [88]: 1 df_shop.iloc[:2].loc[:, 'lastname']
```

```
Out[88]: 0    PERKINS
1    ROBINSON
Name: lastname, dtype: object
```

```
In [89]: 1 df_shop.loc[:, 'lastname'].iloc[:2].str.lower()
```

```
Out[89]: 0    perkins
1    robinson
Name: lastname, dtype: object
```

String Functions in Pandas

```
In [88]: 1 df_shop.iloc[:2].loc[:, 'lastname']
```

```
Out[88]: 0    PERKINS
1    ROBINSON
Name: lastname, dtype: object
```

```
In [89]: 1 df_shop.loc[:, 'lastname'].iloc[:2].str.lower()
```

```
Out[89]: 0    perkins
1    robinson
Name: lastname, dtype: object
```

```
In [90]: 1 df_shop.lastname[:2].str.capitalize()
```

```
Out[90]: 0    Perkins
1    Robinson
Name: lastname, dtype: object
```

String Functions in Pandas

```
In [88]: 1 df_shop.iloc[:2].loc[:, 'lastname']
```

```
Out[88]: 0    PERKINS
          1    ROBINSON
          Name: lastname, dtype: object
```

```
In [89]: 1 df_shop.loc[:, 'lastname'].iloc[:2].str.lower()
```

```
Out[89]: 0    perkins
          1    robinson
          Name: lastname, dtype: object
```

```
In [90]: 1 df_shop.lastname[:2].str.capitalize()
```

```
Out[90]: 0    Perkins
          1    Robinson
          Name: lastname, dtype: object
```

```
In [91]: 1 df_shop.lastname[:2].str.startswith('ROB') # .endswith() , .contains()
```

```
Out[91]: 0    False
          1    True
          Name: lastname, dtype: bool
```

String Functions in Pandas

```
In [88]: 1 df_shop.iloc[:2].loc[:, 'lastname']
```

```
Out[88]: 0    PERKINS
          1    ROBINSON
          Name: lastname, dtype: object
```

```
In [89]: 1 df_shop.loc[:, 'lastname'].iloc[:2].str.lower()
```

```
Out[89]: 0    perkins
          1    robinson
          Name: lastname, dtype: object
```

```
In [90]: 1 df_shop.lastname[:2].str.capitalize()
```

```
Out[90]: 0    Perkins
          1    Robinson
          Name: lastname, dtype: object
```

```
In [91]: 1 df_shop.lastname[:2].str.startswith('ROB') # .endswith() , .contains()
```

```
Out[91]: 0    False
          1    True
          Name: lastname, dtype: bool
```

```
In [92]: 1 df_shop.lastname[:2].str.replace('P', '*')
```

```
Out[92]: 0    *ERKINS
          1    ROBINSON
          Name: lastname, dtype: object
```

String Functions in Pandas

```
In [88]: 1 df_shop.iloc[:2].loc[:, 'lastname']
```

```
Out[88]: 0    PERKINS
          1    ROBINSON
          Name: lastname, dtype: object
```

```
In [89]: 1 df_shop.loc[:, 'lastname'].iloc[:2].str.lower()
```

```
Out[89]: 0    perkins
          1    robinson
          Name: lastname, dtype: object
```

```
In [90]: 1 df_shop.lastname[:2].str.capitalize()
```

```
Out[90]: 0    Perkins
          1    Robinson
          Name: lastname, dtype: object
```

```
In [91]: 1 df_shop.lastname[:2].str.startswith('ROB') # .endswith() , .contains()
```

```
Out[91]: 0    False
          1    True
          Name: lastname, dtype: bool
```

```
In [92]: 1 df_shop.lastname[:2].str.replace('P', '*')
```

```
Out[92]: 0    *ERKINS
          1    ROBINSON
          Name: lastname, dtype: object
```

Pandas datetime functions

Pandas datetime functions

```
In [93]: 1 df_taxi.iloc[:2].tpep_pickup_datetime
```

```
Out[93]: 1    2017-01-05 15:14:52  
2    2017-01-11 14:47:52  
Name: tpep_pickup_datetime, dtype: datetime64[ns]
```

Pandas datetime functions

```
In [93]: 1 df_taxi.iloc[:2].tpep_pickup_datetime
```

```
Out[93]: 1    2017-01-05 15:14:52  
2    2017-01-11 14:47:52  
Name: tpep_pickup_datetime, dtype: datetime64[ns]
```

```
In [94]: 1 df_taxi.iloc[:2].tpep_pickup_datetime.dt.day
```

```
Out[94]: 1      5  
2     11  
Name: tpep_pickup_datetime, dtype: int64
```

Pandas datetime functions

```
In [93]: 1 df_taxi.iloc[:2].tpep_pickup_datetime
```

```
Out[93]: 1    2017-01-05 15:14:52  
2    2017-01-11 14:47:52  
Name: tpep_pickup_datetime, dtype: datetime64[ns]
```

```
In [94]: 1 df_taxi.iloc[:2].tpep_pickup_datetime.dt.day
```

```
Out[94]: 1      5  
2     11  
Name: tpep_pickup_datetime, dtype: int64
```

```
In [95]: 1 df_taxi.iloc[:2].tpep_pickup_datetime.dt.day_of_week
```

```
Out[95]: 1      3  
2      2  
Name: tpep_pickup_datetime, dtype: int64
```

Pandas datetime functions

```
In [93]: 1 df_taxi.iloc[:2].tpep_pickup_datetime
```

```
Out[93]: 1    2017-01-05 15:14:52  
2    2017-01-11 14:47:52  
Name: tpep_pickup_datetime, dtype: datetime64[ns]
```

```
In [94]: 1 df_taxi.iloc[:2].tpep_pickup_datetime.dt.day
```

```
Out[94]: 1      5  
2     11  
Name: tpep_pickup_datetime, dtype: int64
```


```
In [95]: 1 df_taxi.iloc[:2].tpep_pickup_datetime.dt.day_of_week
```

```
Out[95]: 1      3  
2      2  
Name: tpep_pickup_datetime, dtype: int64
```

and more: https://pandas.pydata.org/pandas-docs/stable/user_guide/timeseries.html#time-date-components

Transforming with Train/Test Split

- When performing data transformation



Next Time

- Dimensionality Reduction
 - Feature Selection
 - Feature Extraction

Questions?